

Deep Hierarchical Reinforcement Learning Algorithms in Partially Observable Markov Decision Processes

Ph.D. Dissertation Defense

The world is
a global village and
all the children of
the world are one human family.
May we strive for peace

Presenter: Le Pham Tuyen

Advisor: Prof. TaeChoong Chung, Ph.D.

Artificial Intelligence Laboratory

Department of Computer Science and Engineering

Kyung Hee University, 14th November 2018

Thesis Outline

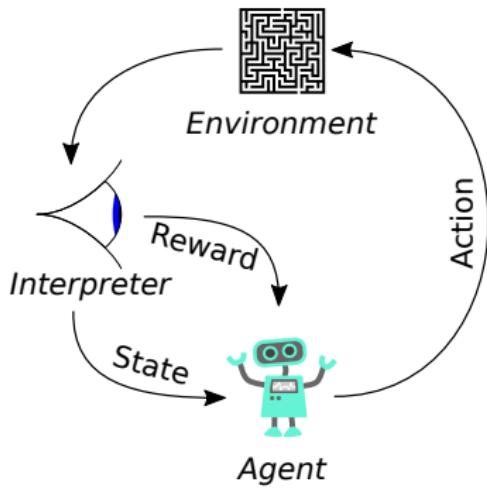
- ① Introduction
- ② Challenges
- ③ Thesis Contributions
- ④ Background and Related Work
- ⑤ Proposed Methodologies
- ⑥ Experiments and Results
- ⑦ Conclusion and Future Work
- ⑧ References

Introduction

Reinforcement Learning

Reinforcement Learning

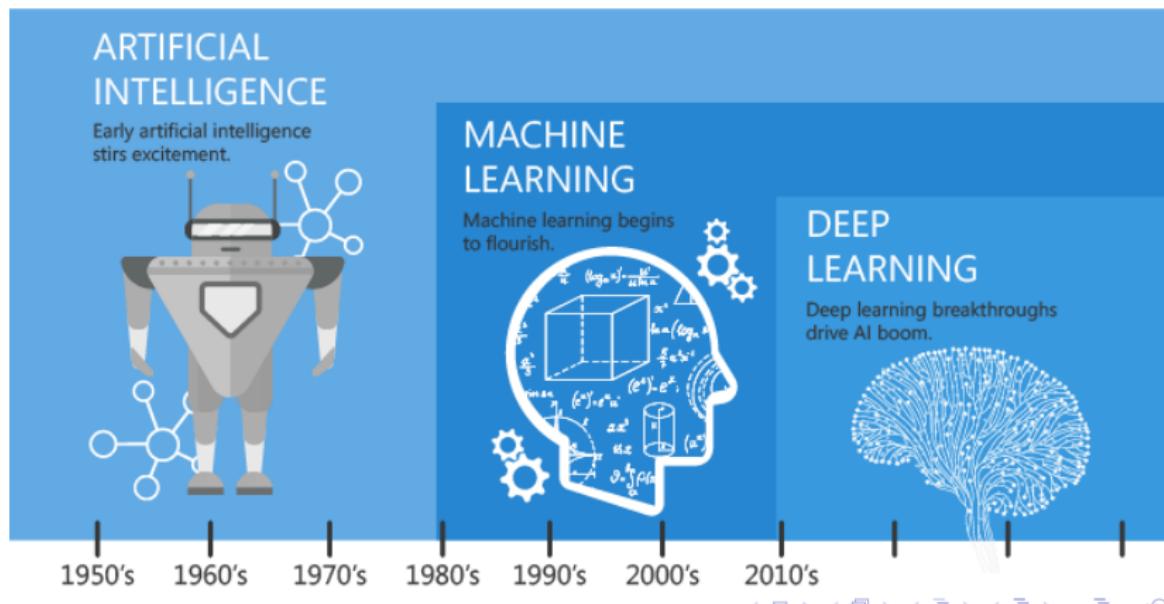
An area of **Machine Learning** concerned with how software agents take actions in an environment so as to maximize cumulative reward.



Machine Learning

- We can answer the 4 major questions:

- ▶ How much/How many?
- ▶ Which category?
- ▶ Which group?
- ▶ Which action?



How much / How many?

- What will be the temperature tomorrow?
- What will be my energy costs next week?
- How many new user will visit next month?

⇒ Regression



Which category?

- Is there a cat or a dog on the image?
- Which emails are spam emails?
- What is the category of this news article (finance, weather, entertainment, sport, . . .)?

⇒ Classification



Which group?

- Which customers have the same favorite product?
- Which visitors like the same movie?
- Which documents has the same topic?

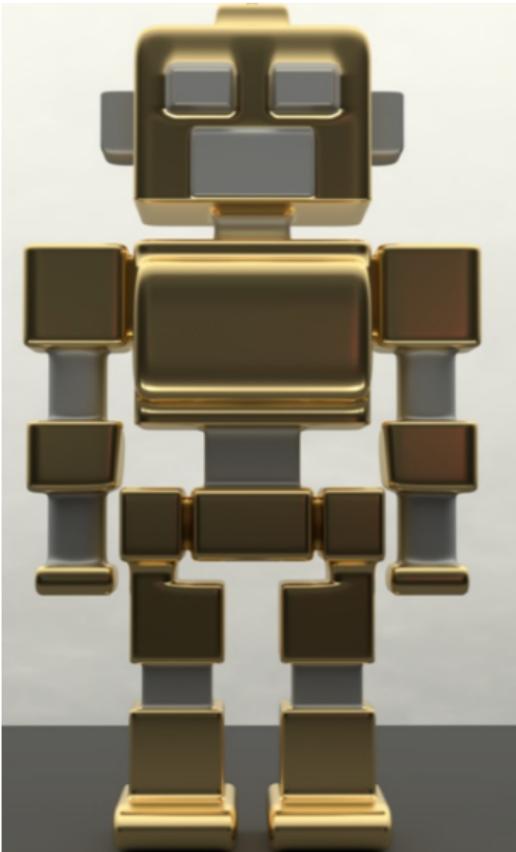
⇒ Clustering



Which action?

- Should I rise or lower the temperature?
- Should I break or accelerate?
- What is the next move for this Go match?

⇒ Reinforcement Learning (RL)



RL application areas

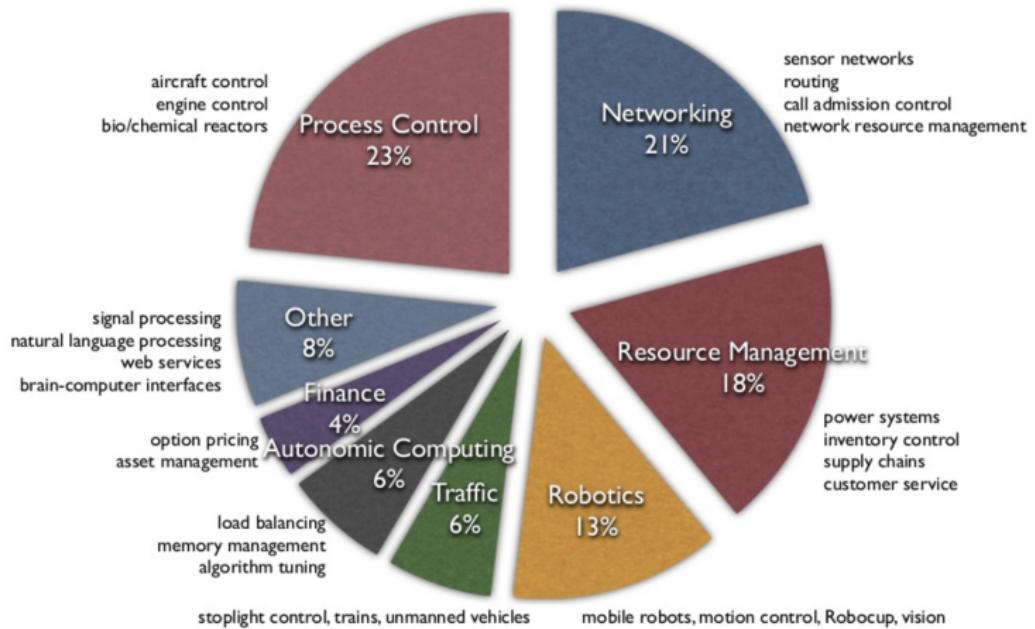


Figure: Rich Sutton. Deconstructing Reinforcement Learning. ICML 09

Era of Deep Reinforcement Learning

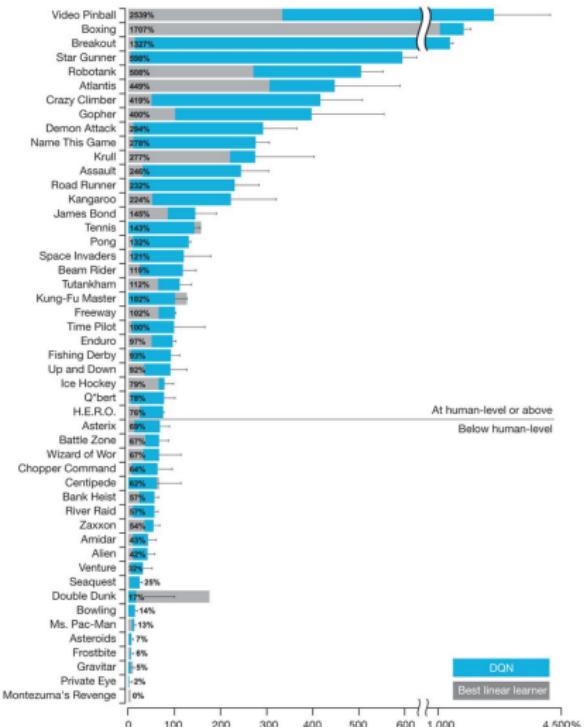


Figure: DQN in Atari Games



(a) Go game



(b) Starcraft



(c) DotA

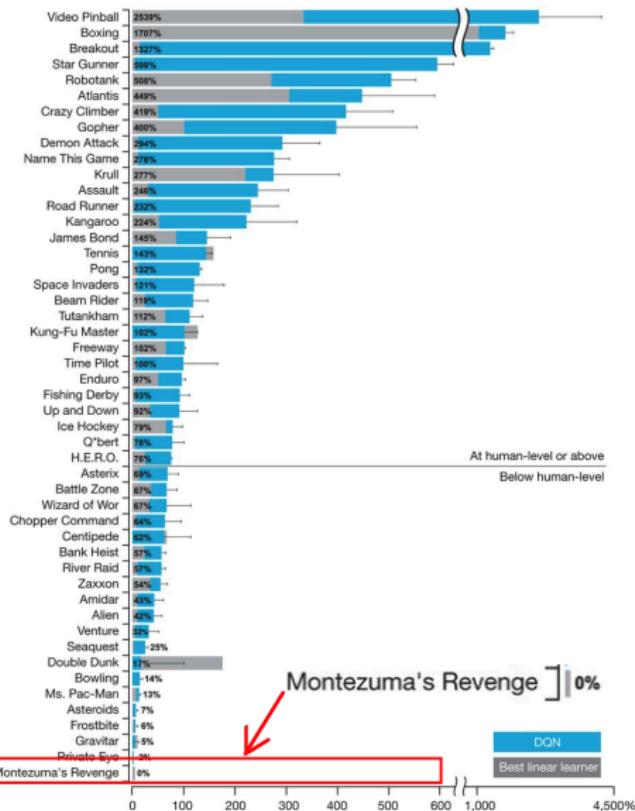


(d) Poker

Figure: Domains which the agent defeats human

Challenges

Challenge 1



Hierarchical Task

*DQN as well as plain DRL algorithms fails to solve the task having multiple subtasks (**hierarchical task**) such as Montezuma's Revenge in Atari Game 2600*



Montezuma's Revenge Game

Challenge 2

Partial Observability

- Most of studies assume that an agent can observe the environment states fully (**MDP**)
- However, it does not reflect the nature of real-world applications, where the agent only observes a partial states (**POMDP**)

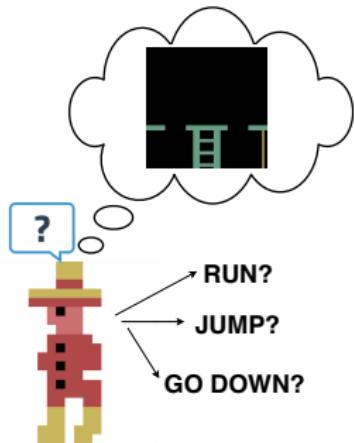
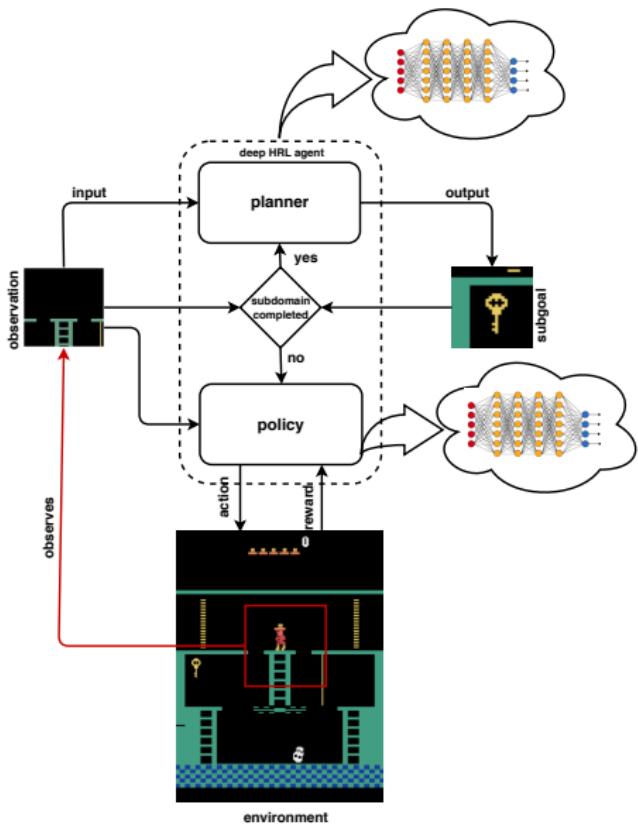


Figure: The agent takes the action under partial observability

Proposed Concept



We want to propose a deep HRL algorithm for solving **hierarchical tasks** under **partial observability**

- The proposed frameworks employ deep neural network as policies.
- The proposed frameworks use limited observations to make decisions.
- The proposed frameworks can solve hierarchical tasks

Thesis Contributions

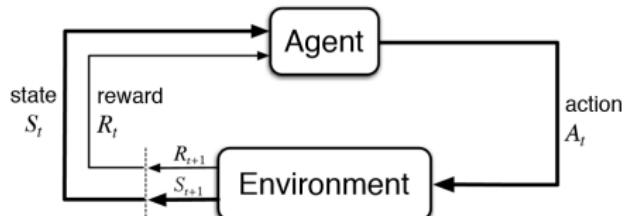
- **Develop:** **hierarchical Deep Recurrent Q-Learning algorithms (hDRQNs)** in order to handle **hierarchical tasks** in **POMDP**. Particularly,
 - ▶ We develop hDRQNv1 algorithm which learns a framework of hierarchical polices.
 - ★ Two levels of hierarchical policies: meta-controller is the upper policy and sub-controller is the lower policy.
 - ★ Two hierarchical policies integrated recurrent neural networks are expected to overcome the challenges under partial observability
 - ▶ We develop hDRQNv2 algorithm of a proposed framework which integrates recurrent neural networks in a different way, thus expected to have better performance.
- To the best of our knowledge, our research is **the first study** that learns Montezuma's Revenge under partial observability.

Background and Related Work

- Reinforcement learning (Markov Decision Process)
- Hierarchical reinforcement learning (Semi Markov Decision Process)
- Planing under partial observability (Partial Observation Markov Decision Process)
- Related works:
 - ▶ Deep Q Networks (DQN)
 - ▶ Deep Recurrent Q Network (DRQN)
 - ▶ Hierarchical Deep Q Network (hDQN)

Markov Decision Process (MDP)

- RL can be formalized as a MDP with five elements $\{\mathcal{S}; \mathcal{A}; r; \mathcal{P}; \gamma\}$



- ▶ \mathcal{S} state space
- ▶ \mathcal{A} action space
- ▶ $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ reward function
- ▶ $\mathcal{P}(s'|s, a)$ transition dynamics
- ▶ γ discount factor

- Markov property: $\mathcal{P}(s_{t+1}|s_1, a_1, \dots, s_t, a_t) = \mathcal{P}(s_{t+1}|s_t, a_t)$
- A policy π is a map from state to action. E.g.
 - ▶ Deterministic policy: $a = \pi(s)$
 - ▶ Stochastic policy: $\pi(a|s) = P[a_t = a | s_t = s]$

Goal of RL

Find an optimal policy π^* in order to maximize the expected discounted reward: $J(\pi) = \mathbb{E} \left[\sum_{t=1}^{\infty} \gamma^{t-1} r(a_t, s_t) \right]$

Partial Observation

Markov Decision Process (POMDP)

- Agent observes the entire environment → **MDP**
- Agent only observes a part of environment → **POMDP**
- **POMDP** is popular in the real-world applications. E.g.
 - ▶ A robot with camera vision isn't told its absolute location
 - ▶ A trading agent only observes current prices
 - ▶ A poker playing agent only observes public cards



(a) Robot Navigation



(b) Trading Bot



(c) Poker Bot

Some POMDP domains

Partial Observation

Markov Decision Process (POMDP)

- POMDP is defined as a tuple of six components $\{\mathcal{S}, \mathcal{A}, \mathcal{P}, r, \Omega, \mathcal{Z}\}$
 - ▶ $\mathcal{S}, \mathcal{A}, \mathcal{P}, r$ are the state space, action space, transition function and reward function, respectively, as in a MDP.
 - ▶ Ω and \mathcal{Z} are the observation space and observation model, respectively
- The agent cannot observe the whole environment, thus, maintain a hidden state b called *belief state*

Definition

Belief state defines the probability of being in state s according to its history of actions and observations; and can be updated using the Bayes rule:

$$b'(s') \propto \mathcal{Z}(o|s', a) \sum_{s \in \mathcal{S}} \mathcal{P}(s'|s, a) b(s).$$

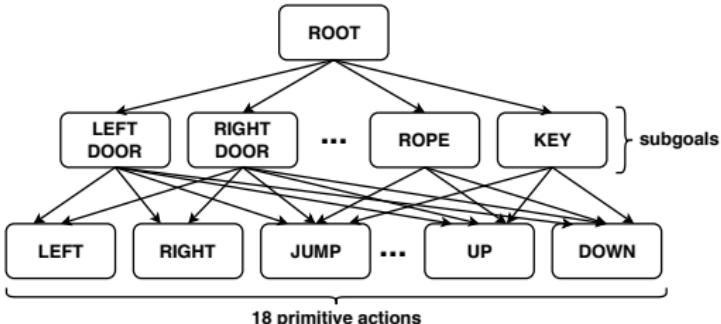
- Updating belief state require a high computational cost and expensive memory → take advantages of RNNs

Semi Markov Decision Process (SMDP)

- Hierarchical tasks are popular in real-world applications. E.g.
 - An agent navigates to the key before reaching the door to open.
 - Tasks of a taxi: go to to the passengers, pick up, go to to the destination, take off.
 - A robot plans to go to the door before going to the destination.



(a) Montezuma's Revenge

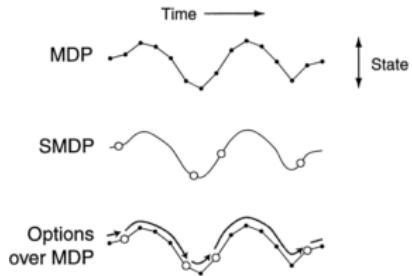


(b) The hierarchy of Montezuma's Revenge domain

Hierarchical Domain

- SMDP** is an extensional theory of MDP, was developed to deal with challenges in hierarchical tasks.

- SMDP = Options over MDP.

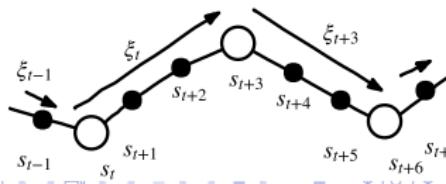


Definition

An *option* $\xi \in \Xi$ is defined by three elements:

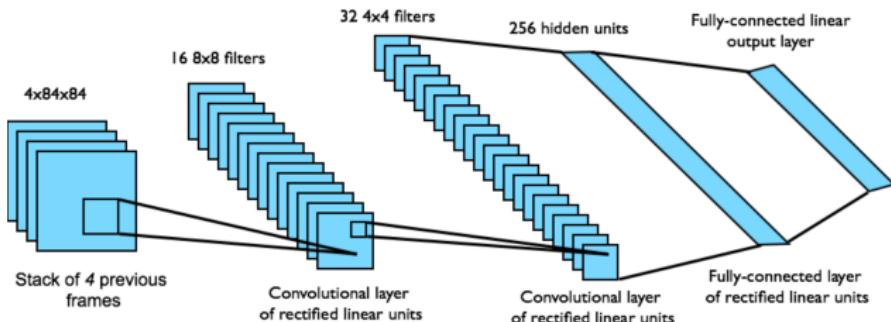
- An option's policy π ,
- A termination condition β
- An initiation set $\mathcal{I} \subseteq \mathcal{S}$ denoted as the set of states in the option

- A *policy over options* $\mu(\xi|s)$ is introduced to select options
- An option is executed as follows:
 - ▶ Under option ξ_t , state s_t , the action a_t is selected based on π
 - ▶ The environment transits to state s_{t+1}
 - ▶ The option executes until state s_{t+3}
 - ▶ The next option is selected $\xi_{t+3} = \mu(s_{t+3})$



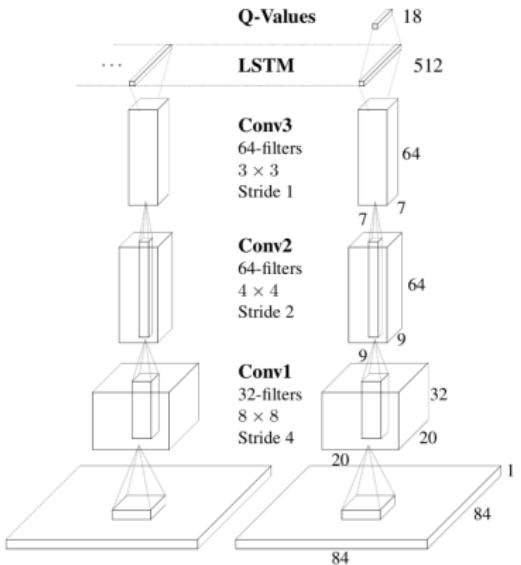
Deep Reinforcement Learning (1)

- Deep Q Learning (DQN) for Atari Games
 - ▶ End-to-end learning of values $Q(s, a)$ from raw pixels
 - ▶ Input state s is stack of raw pixels from last 4 frames
 - ▶ Output is $Q(s, a)$ for 18 joystick/button positions
 - ▶ Hidden layers are the combination of CONV, FC, ReLU
 - ▶ Stabilization techniques:
 - ★ Experience replay.
 - ★ Delayed target network.



- Other tricks:
 - ▶ Double Deep Q Learning (DDQN)
 - ▶ Dueling network
 - ▶ Prioritized replay

Deep Reinforcement Learning (2)

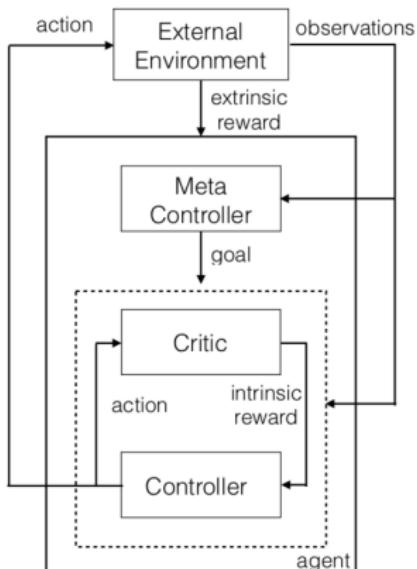


- Limitations of DQN and its derivations:
 - ▶ Only learning from a limited number of past states (last 4 frames)
 - ▶ Cannot deal with POMDP domains
- Deep Recurrent Q-Network(DRQN) [6]:
 - ▶ A combination of a Long Short Term Memory (LSTM) and a DQN
 - ▶ Better handles the loss of information (POMDP)

- Other tricks combining with DRQN [6]:

- ▶ Updating DRQN techniques: *Bootstrapped Sequential Updates* vs *Bootstrapped Random Updates*
- ▶ Ignore first observations in a sequence of transitions when updating the Q value function

Deep Reinforcement Learning (3)



- hDQN framework [3]

- ▶ Two levels of controllers: *meta controller* and *controller*
- ▶ The *meta controller* produces a subgoal.
- ▶ The *controller* performs primitive actions to obtain the subgoal.
- ▶ The set of subgoals is predefined and fixed.
- ▶ The *meta controller* and the *controller* are built from DQN networks
- ▶ *Extrinsic* is reward of the meta controller and *intrinsic* is reward of the controller
- ▶ Only deal with fully observable domains

- Others:

- ▶ Option Critic framework [1] and Feudral framework [2]
- ▶ Discovering subgoals [4]
- ▶ Adaptively finding a number of options [5]

Proposed Methodologies

hDRQN: Key Terminologies (1)

Subdomain (ξ)

- A domain = multiple subdomains.
- A subdomain \Leftrightarrow an option ξ .

E.g. Domain: Montezuma's Revenge.

Subdomains: move to the left door, move to the key, ...

Subgoal (g)

Each subdomain has a subgoal $g \in \Omega$

E.g. White rectangles (left image)

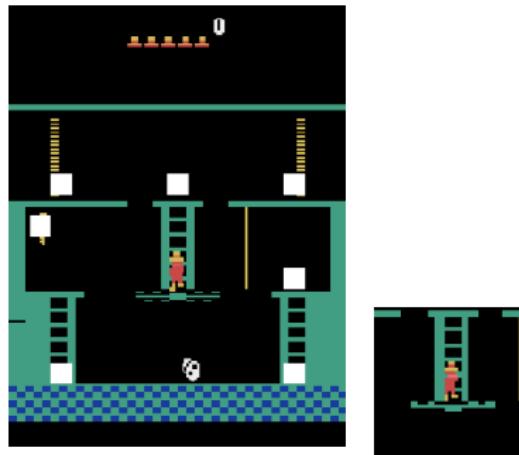


Figure: Montezuma's Revenge

Observation (o)

A partial of the environment ($o \in \Omega$) which the agent can observe

E.g. The pixels around the agent (right image)

hDRQN: Key Terminologies (2)

Meta-controller (META)

*Equivalent to a “**policy over subgoals**” that receives the current observation o_t and determines the new subgoal g_t*

- E.g. In Montezuma's Revenge, META is used to select new subgoal.

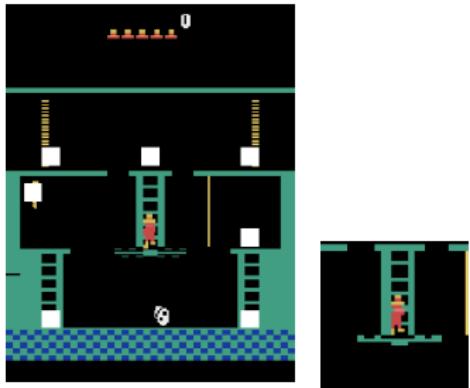


Figure: Montezuma's Revenge

Extrinsic Reward (r^{ex})

Use to evaluate the goodness of META.

- E.g. In Montezuma's Revenge, $r^{ex} = 1$ if agent obtains the key or opens the doors, otherwise 0

hDRQN: Key Terminologies (3)

Sub-controller (SUB)

*Equivalent to the **option's policy**, which directly interacts with the environment by performing action a_t*

- E.g. In Montezuma's Revenge, SUB controls the agent to move between subgoals.

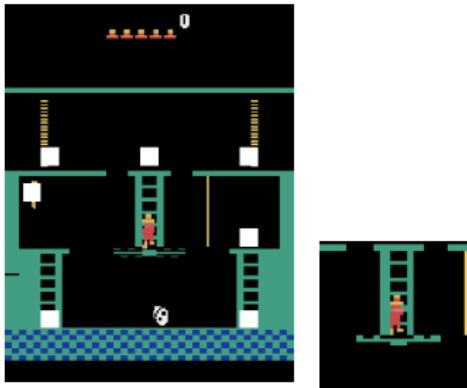


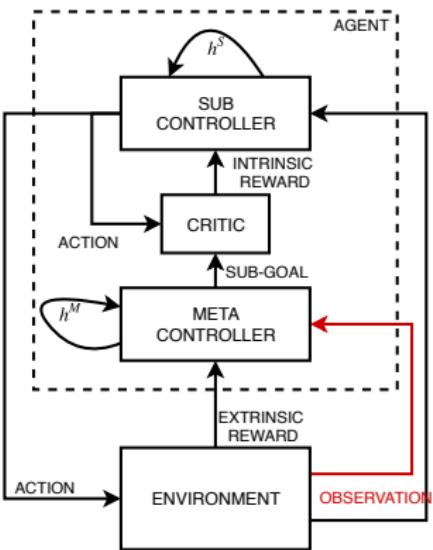
Figure: Montezuma's Revenge

Intrinsic Reward (r^{in})

Use to evaluate the goodness of SUB.

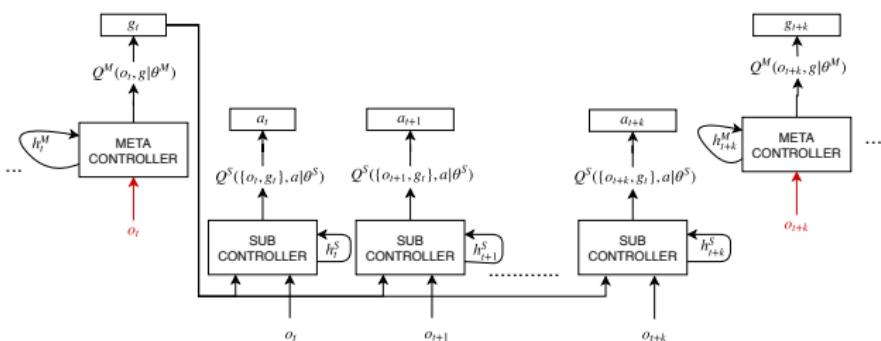
- E.g. In Montezuma's Revenge, $r^{in} = 1$ if agent obtains the subgoal, otherwise 0

hDRQN: Framework 1



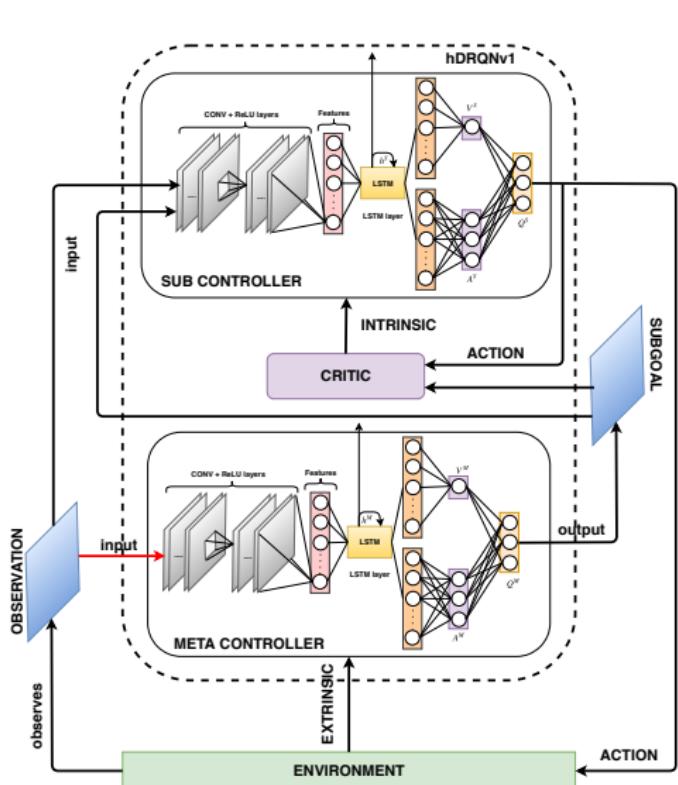
• hDRQNv1:

- ▶ Inspired by hDQN framework [3]
- ▶ Build on two deep **recurrent** neural policies.
- ▶ Input is a single frame (hDQN uses 4 frames)



hDRQN: Framework 1 (Extended)

META:

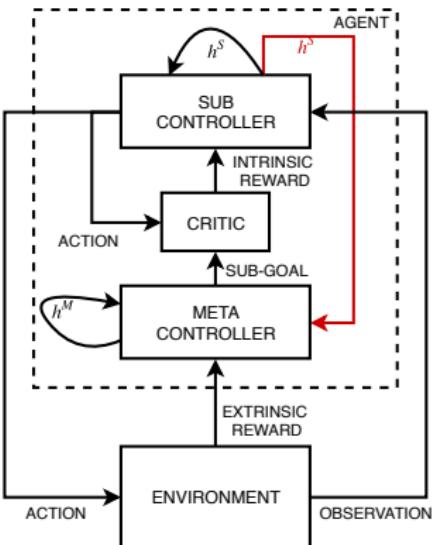


- Input: Observation o
- Feature extraction: 4 CONV layers and ReLU layers.
- LSTM is integrated in front of the features.
- The output of LSTM is put into Dueling network ([7])
- Output: Q subgoal values $Q^M(o, g)$

SUB:

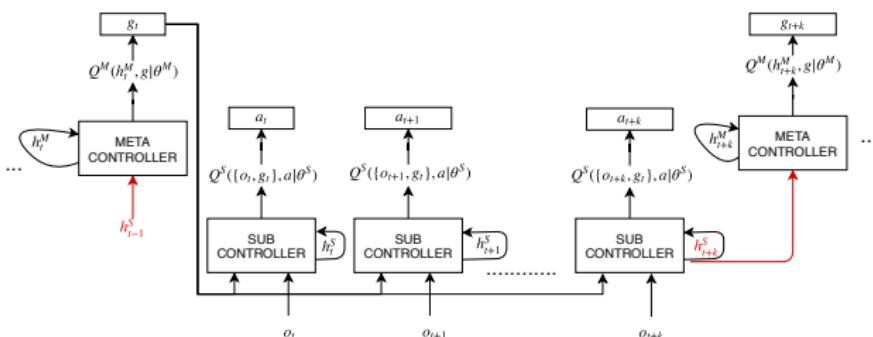
- Input: Observation o and current subgoal (g)
- Other part: same as META
- Output: Q action values $Q^S(\{o, g\}, a)$

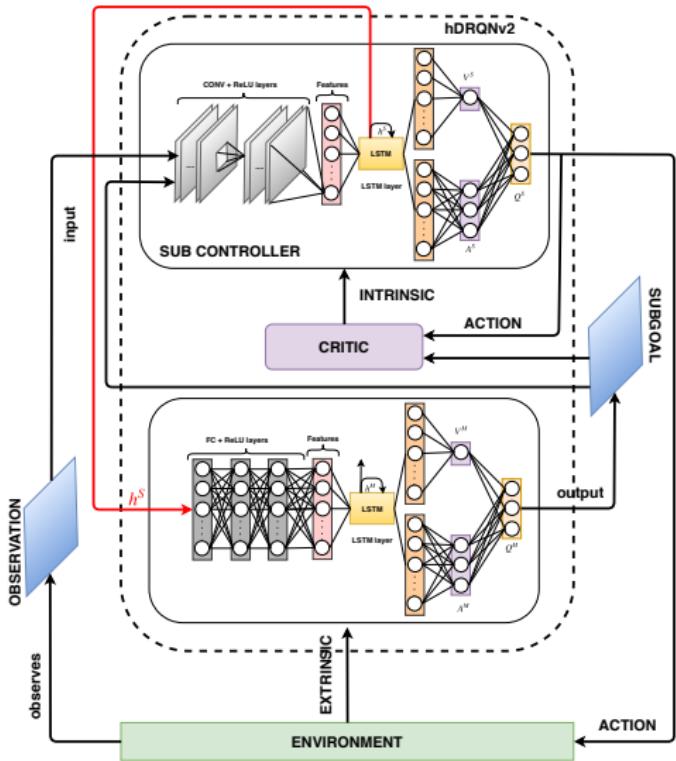
hDRQN: Framework 2



• hDRQNv2

- ▶ An improved version of hDRQNv1
- ▶ Input of META is the internal states of LSTM layer in SUB





META:

- Input: hidden states from SUB h^S
- Feature extraction: Three fully connected layers and ReLU layers.
- Other part has the same architecture as META of framework 1

SUB:

- Same architecture as SUB of framework 1

hDRQN: Q values

- META Q subgoal values:

$$h_t^M, Q^M(o_t, g_t) = f^M(\Phi^M, h_{t-1}^M)$$

- SUB Q action values:

$$h_t^S, Q^S(\{o_t, g_t\}, a_t) = f^S(\Phi^S, h_{t-1}^S)$$

- Where:

- f^M and f^S are the recurrent networks of the META and SUB.
- h_t^M and h_t^S are internal states constructed by recurrent networks.
- Φ^M and Φ^S are the features of META and SUB.

$$\Phi^M = \begin{cases} f^{extract}(o_t) & \text{framework 1} \\ f^{extract}(h_t^S) & \text{framework 2} \end{cases}$$

$$\Phi^S = f^{extract}(o_t, g_t)$$

- $f^{extract}$ is neural networks to extract features from input (E.g. CONV, FC, ReLU, ...)

hDRQN: Learning META

- Optimizing META by minimizing loss functions:

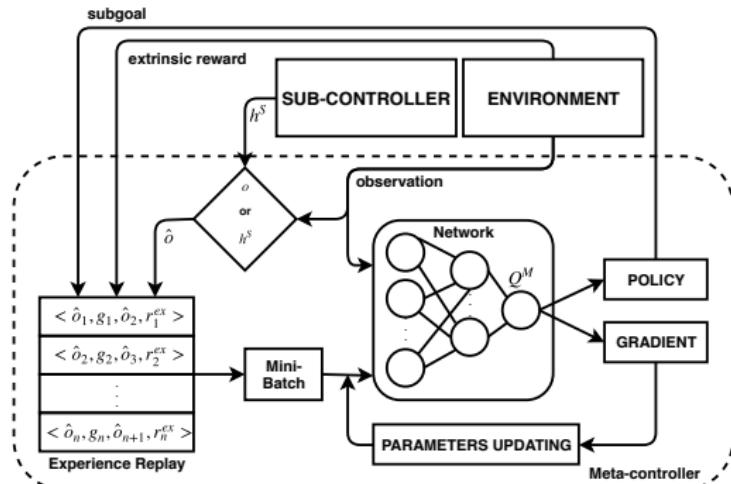
$$\mathcal{L}^M = \mathbb{E}_{(o, g, o', g', r^{ex}) \sim \mathcal{M}^M} [y_i^M - Q^M(o, g)]$$

- Where:

► y_i^M is target values of META

$$y_i^M = r^{ex} + \gamma Q^{M'}(o', \operatorname{argmax}_{g'} Q^M(o', g'))$$

- Minibatch Sampling Strategy: Bootstrapped Random Updates [6].



hDRQN: Learning SUB

- Optimizing SUB by minimizing loss functions:

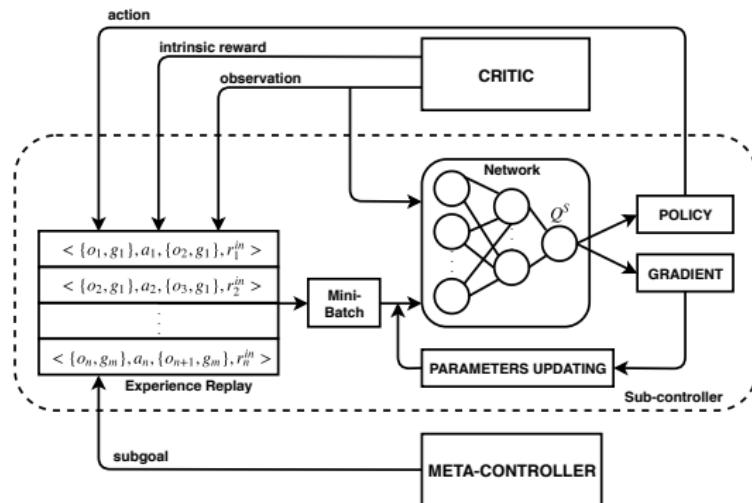
$$\mathcal{L}^S = \mathbb{E}_{(o,g,a,r^{in}) \sim \mathcal{MS}} [y_i^S - Q^S(\{o, g\}, a)]$$

- Where:

► y_i^S are target values of SUB

$$y_i^S = r^{in} + \gamma Q^{S'}(\{o', g\}, \text{argmax}_{a'} Q^S(\{o', g\}, a'))$$

- Minibatch Sampling Strategy: Bootstrapped Random Updates [6].



hDRQN: Sampling Strategy

- Bootstrapped Random Updates [6] is compatible with recurrent neural networks:
 - Randomly selects a batch of episodes from the experience replay
 - For each episode, we begin at a random transition and select a sequence of n transitions
 - For each controller, we have n^M (META) and n^S (SUB)

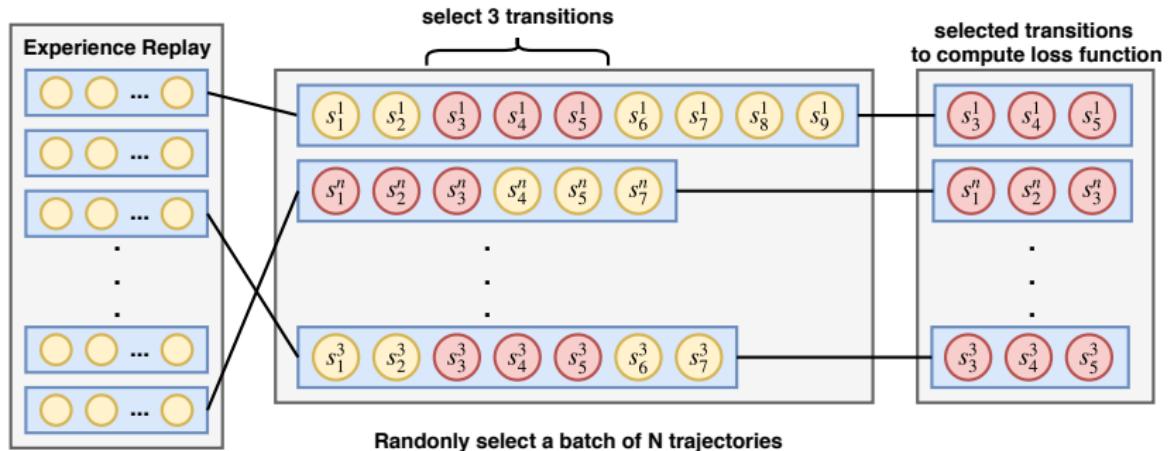


Figure: Bootstrapped Random Updates

hDRQN: Pseudo code

Algorithm 1 hDRQN in POMDP

Require:

- 1: POMDP $M = \{\mathcal{S}, \mathcal{A}, \mathcal{P}, r, \Omega, \mathcal{Z}\}$
- 2: Meta-controller with the network Q^M (main) and $Q^{M'}$ (target) parameterized by θ^M and $\theta^{M'}$, respectively.
- 3: Sub-controller with the network Q^S (main) and $Q^{S'}$ (target) parameterized by θ^S and $\theta^{S'}$, respectively.
- 4: Exploration rate ϵ^M for meta-controller and ϵ^S for sub-controller.
- 5: Experience replay memories M^M and M^S of meta-controller and sub-controller, respectively.
- 6: A pre-defined set of subgoals \mathcal{G} .
- 7: f^M and f^S are recurrent networks of meta-controller and sub-controller, respectively.

Ensure:

- 8: **Initialize:**
 - Experiences replay memories M^M and M^S
 - Randomly initialize θ^M and θ^S
 - Assign value to the target networks $\theta^{M'} \leftarrow \theta^M$ and $\theta^{S'} \leftarrow \theta^S$
 - $\epsilon^M \leftarrow 1.0$ and decreasing to 0.1
 - $\epsilon^S \leftarrow 1.0$ and decreasing to 0.1
- 9: **for** $k = 1, 2, \dots, K$ **do**
- 10: **Initialize:** the environment and get the start observation o

- 11: **Initialize:** hidden states $h^M \leftarrow \mathbf{0}$
- 12: **while** o is not terminal **do**
- 13: **Initialize:** hidden states $h^S \leftarrow \mathbf{0}$
- 14: **Initialize:** start observations $o_0 \leftarrow \hat{o}$ where \hat{o} car be observation o or hidden state h^S
- 15: **Determine subgoal:** $g, h^M \leftarrow EPS_GREEDY(\hat{o}, h^M, \mathcal{G}, \epsilon^M, Q^M, f^M)$
- 16: **while** o is not terminal and g is not reached **do**
- 17: **Determine action:** $a, h^S \leftarrow EPS_GREEDY(\{o, g\}, h^S, \mathcal{A}, \epsilon^S, Q^S, f^S)$
- 18: **Execute** action a , receive reward r , extrinsic reward r^{ex} , intrinsic reward r^{in} , and obtain the next state s'
- 19: **Store transition** $\{(o, g), a, r^{in}, \{o', g'\}\}$ in M^S
- 20: **Update sub-controller**
 $SUB_UPDATE(M^S, Q^S, Q^{S'})$
- 21: **Update meta-controller**
 $META_UPDATE(M^M, Q^M, Q^{M'})$
- 22: **Transition to next observation** $o \leftarrow o'$
- 23: **end while**
- 24: **Store transition** $\{o_0, g, r_{total}^{ex}, \hat{o}'\}$ in M^S where \hat{o} can be observation o' or the last hidden state h^S
- 25: **end while**
- 26: **Anneal** ϵ^M and ϵ^S
- 27: **end for**

Experiments and Results

Experiments

- Domains:

- ▶ Multiple goals in gridworld.
- ▶ Multiple goals in four-rooms.
- ▶ Montezuma's Revenge.

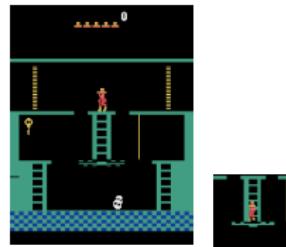
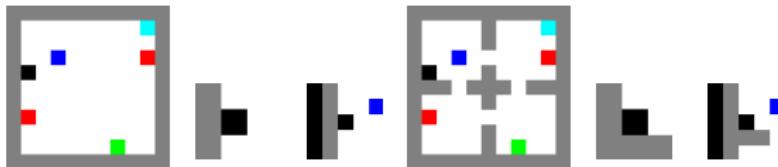


Figure: Domains

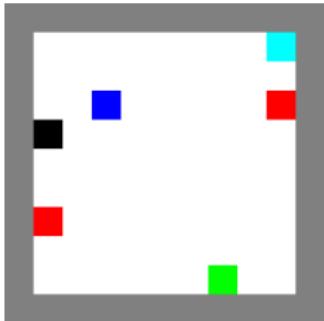
- Implementation details:

- ▶ Tensorflow.
- ▶ The inputs of META and SUB are a raw image of size $44 \times 44 \times 3$
- ▶ Feature size is 256
- ▶ Input and output of LSTM have 256 values.
- ▶ Using ADAM to optimize the controller's parameters
- ▶ Learning rate is 0.001
- ▶ Discount factor is 0.99

Domain Description (1)

- Multiple goal in Gridworld:

- Gridworld map of size 11×11 .
- 4 types of objects: an agent (in black), two obstacles (in red) and two goals (in blue and green) or three goals (in blue, green and cyan)
- Objects are randomly located on the map
- Four actions: top, down, left or right.



- Reward:

- Proper order: blue \Rightarrow green (two goals) or blue \Rightarrow green \Rightarrow cyan (three goals)

- Classical reward:

$$r = \begin{cases} 1 & \text{for each reached goals in proper order} \\ -1 & \text{hit the obstacle} \end{cases}$$

- Intrinsic reward:

$$r^{in} = \begin{cases} 1 & \text{obtain the goal} \\ -1 & \text{hit the obstacle} \end{cases}$$

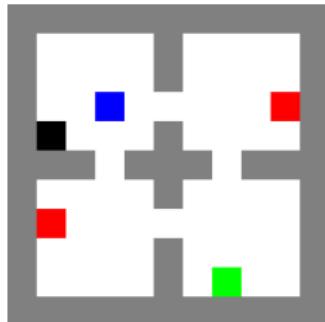
- Extrinsic reward:

$$r^{ex} = \begin{cases} 1 & \text{for each reached goal in proper order} \\ 0.01 & \text{otherwise} \end{cases}$$

Domain Description (1)

- Multiple goal in Four-rooms:

- Four-rooms map of size 11×11 .
- 4 types of objects: an agent (in black), two obstacles (in red) and two goals (in blue and green) or three goals (in blue, green and cyan)
- Objects are randomly located on the map
- Four actions: top, down, left or right.



- Reward:

- Proper order: blue \Rightarrow green (two goals) or blue \Rightarrow green \Rightarrow cyan (three goal)

- Classical reward:

$$r = \begin{cases} 1 & \text{reach goals in proper order} \\ -1 & \text{hit the obstacle} \end{cases}$$

- Intrinsic reward:

$$r^{in} = \begin{cases} 1 & \text{obtain the goal} \\ -1 & \text{hit the obstacle} \end{cases}$$

- Extrinsic reward:

$$r^{ex} = \begin{cases} 1 & \text{reach goals in order} \\ 0.01 & \text{otherwise} \end{cases}$$

Domain Description (1)

- Montezuma's Revenge:

- One of the hardest games in ATARI 2600
- DQN achieved a score of zero
- We use OpenAI Gym to simulate this domain
- To pass through the doors, first, the agent needs to pick up the key.
- Agent observes an area of 70×70 pixels



- Reward:

- Classical reward: The agent will earn 100 points after it obtains the key and 300 after it reaches any door
- Intrinsic reward:

$$r^{in} = \begin{cases} 1 & \text{reach subgoal} \\ 0 & \text{otherwise} \end{cases}$$

- Extrinsic reward:

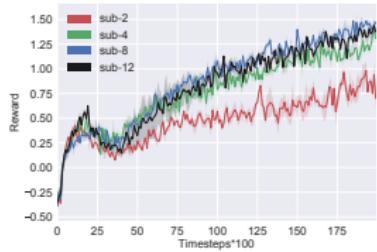
$$r^{ex} = \begin{cases} 1 & \text{obtain key or open door} \\ 0 & \text{otherwise} \end{cases}$$

Experiments

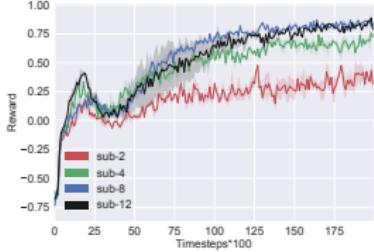
- Experiment 1: Evaluate on different values of n^M and n^S .
 - ▶ Two goals in Grid World
 - ▶ Effect of n^S
 - ▶ Effect of n^M
- Experiment 2: Evaluate on different levels of observation.
 - ▶ Two goals in Grid World
 - ▶ 3×3 observable agent
 - ▶ 5×5 observable agent
 - ▶ Fully observable agent
- Experiment 3: Compare performance of hDRQNv1, hDRQNv2 with:
 - ▶ Flat algorithms (DQN, DRQN)
 - ▶ Hierarchical algorithm (hDQN)
- Experiment 4: Montezuma's Revenge
 - ▶ Successful rate of reaching key
 - ▶ Number of times to visit the subgoals

Experiment 1: Effect of n^S (1)

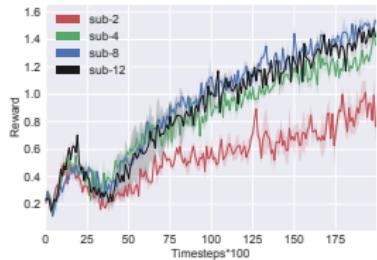
- Report of hDRQNv1 with different n^S (2,4,8,12)



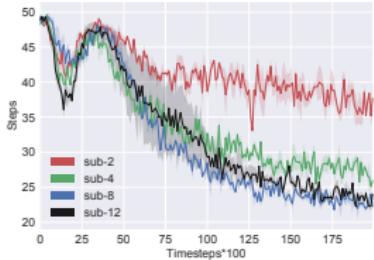
(c) Reward



(d) Intrinsic



(e) Extrinsic

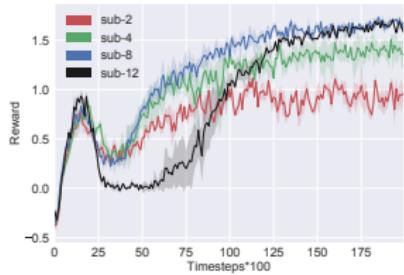


(f) Steps

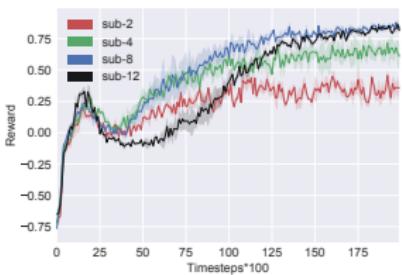
- Fixed $n^M = 1$
- Perform well with a big n^S (8,12)
- Performance decreases when n^S is decreased
- Only a little difference in performance between 8 and 12
- Intuitively, LSTM in SUB needs a long sequence of transitions

Experiment 1: Effect of $n^S(2)$

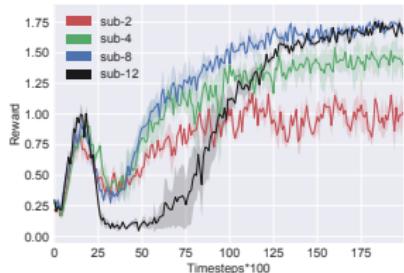
- Report of hDRQNv2 with different n^S (2,4,8,12)



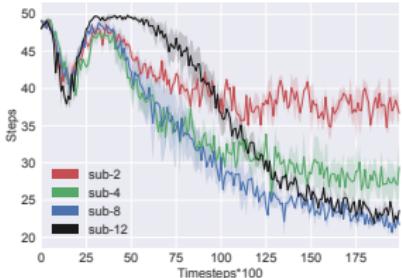
(g) Reward



(h) Intrinsic



(i) Extrinsic

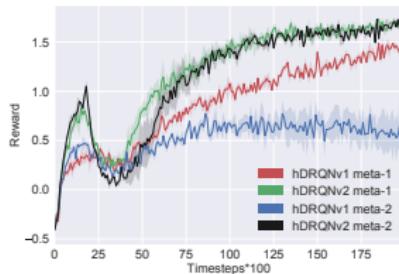


(j) Steps

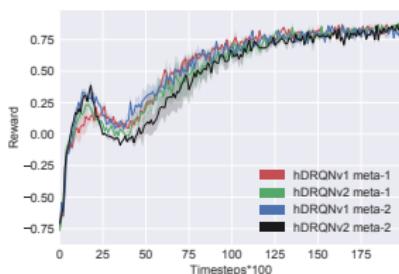
- Fixed $n^M = 1$
- Same behavior as hDRQNv1

Experiment 1: Effect of n^M

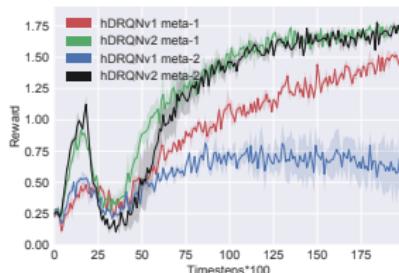
- Report of hDRQNv1 and hDRQNv2 with different n^M (1, 2)



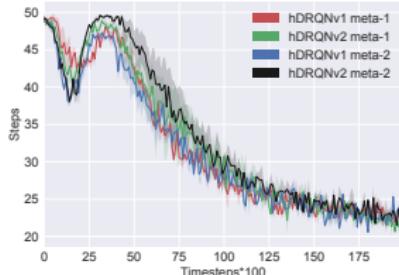
(k) Reward



(l) Intrinsic



(m) Extrinsic

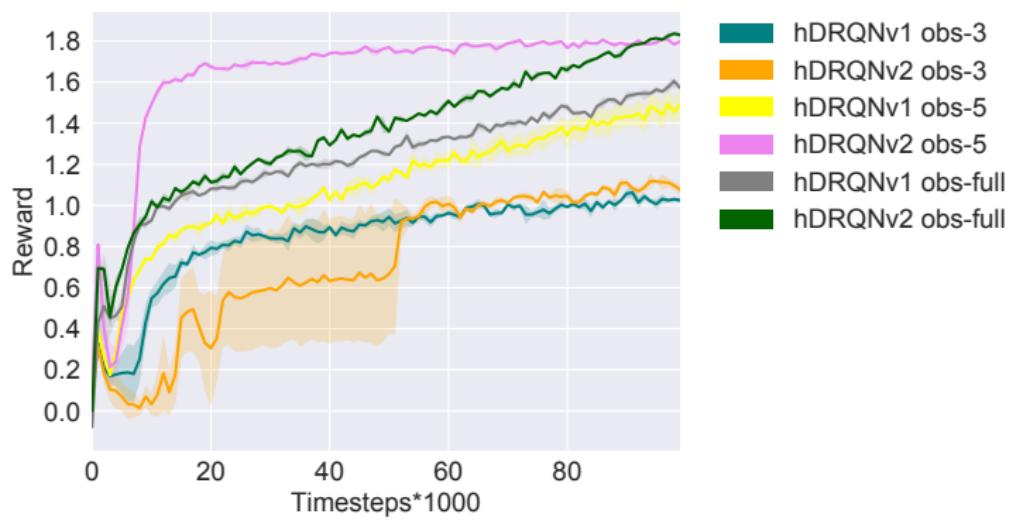


(n) Steps

- Fixed $n^S = 8$
- With hDRQNv1, $n^M = 1$ is better than $n^M = 2$
- With hDRQNv2, the performance is the same at both settings $n^M = 1$ and $n^M = 2$

Experiment 2: Effect of different levels of observation

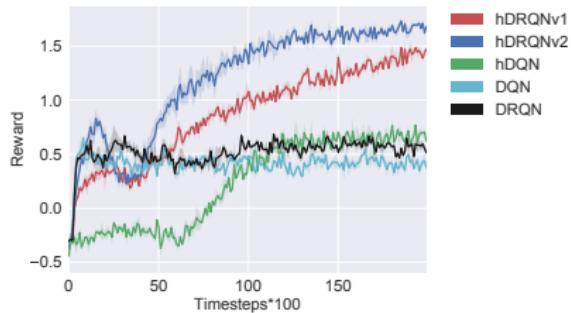
- Performance of the agent with a larger observation area is better than the agents with smaller observing abilities
 - The performance of a 5×5 observable agent using hDRQNv2 seems to converge faster than a fully observable agent



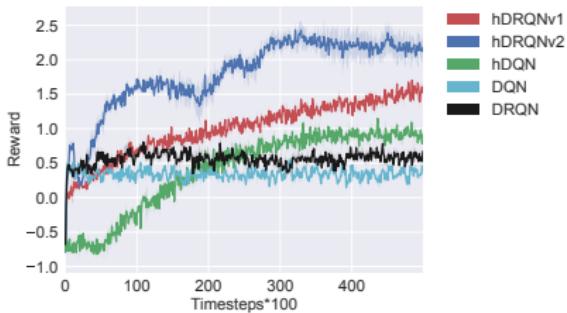
Experiment 3: Performance Comparison (1)

- Multiple goals in gridworld

- The hDRQN algorithms outperforms the other algorithms
- hDRQNV2 has the best performance
- The hDQN algorithm has poor performance in POMDP domains



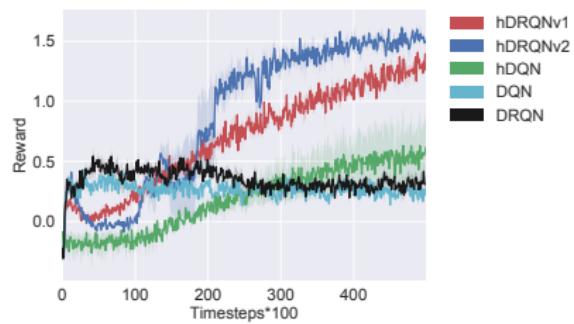
(o) Two goals in Gridworld



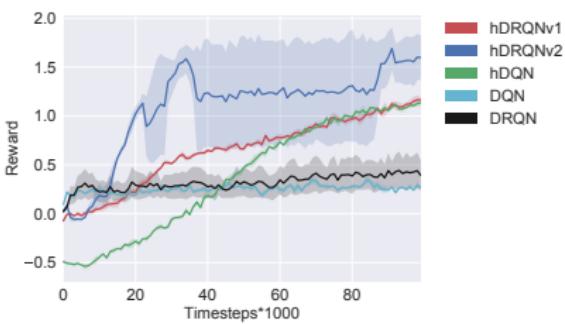
(p) Three goals in Gridworld

Experiment 3: Performance Comparison (2)

- Multiple goals in four-rooms
 - Same behaviors as in Gridworld



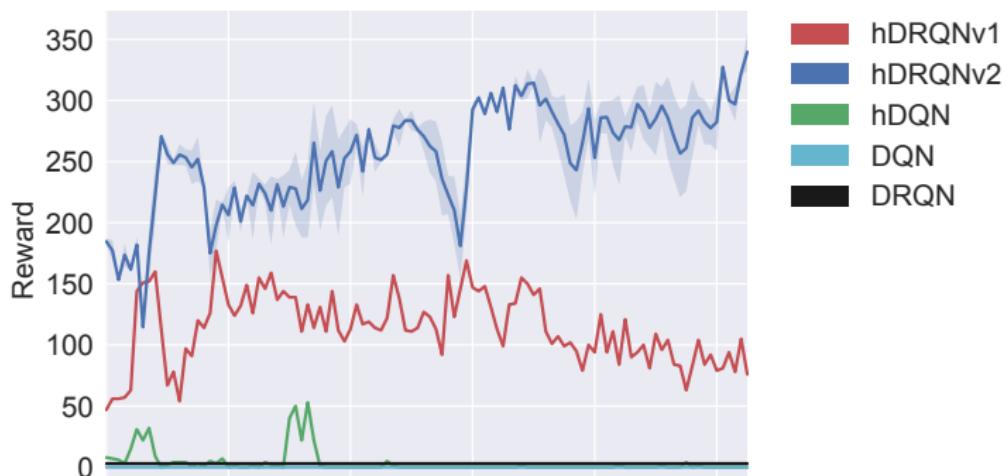
(q) Two goals in Four-rooms



(r) Three goals in Four-rooms

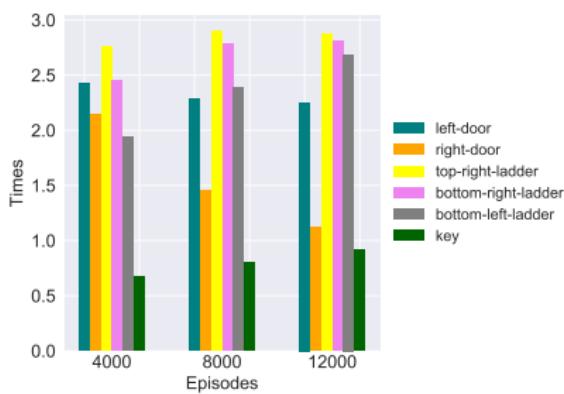
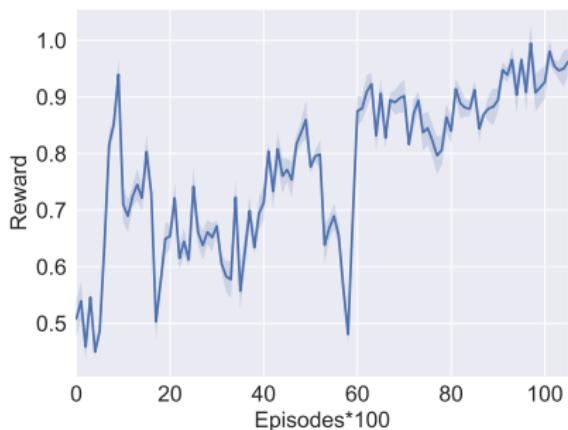
Montezuma's Revenge (1)

- DQN reported a score of zero
- DRQN also achieved a score of zero because of the highly hierarchical complexity of the domain
- hDQN can achieve a high score on this domain
- The hDRQNV2 algorithm shows a better performance than hDRQNV1
⇒ Difference in the architecture of two frameworks has affected their performance



Experiment 4: Montezuma's Revenge (2)

- The agent using the hDRQNv2 algorithm almost picks up the “key” at the end of the learning process
- hDRQNv2 tends to explore more often for subgoals that are on the way to reaching the “key” (E.g. top-right-ladder, bottom-right-ladder, and bottom-left-ladder)
- Exploring less often for other subgoals such as the left door and right door



(s) Success ratio

(t) Number of visits subgoals

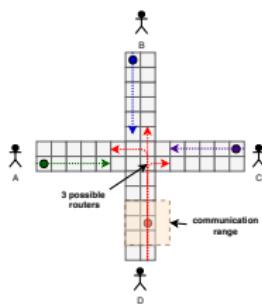
Demo

Conclusions and Future Works

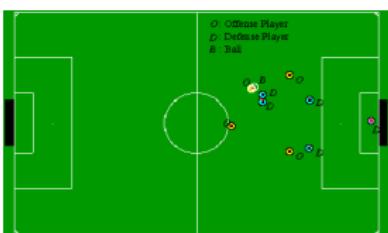
- **Implemented:** new hierarchical deep reinforcement learning algorithms (hDRQNs)
 - ▶ For hierarchical tasks
 - ▶ For both MDP and POMDP tasks
 - ▶ Takes advantage of deep neural networks (DNN, CNN, LSTM)
- **Proposed:** a new way to integrate LSTM into the learning framework, which allows to learning data efficiently and better convergence.
- **Employed:** several advanced methods in deep reinforcement learning:
 - ▶ Double Q Learning
 - ▶ Deep Recurrent Q Network
 - ▶ Dueling Q Network
 - ▶ Bootstrapped Random Updates

Future works

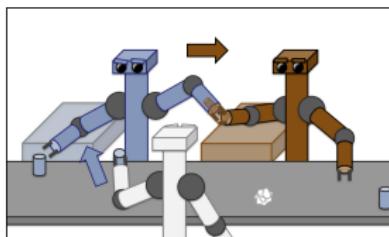
- **Improved:** our framework by tackling those problems:
 - ▶ Our framework is hard to scale for domains with more than two levels of hierarchy
 - ▶ Discovering a set of subgoals in POMDP is still a difficult problem.
- **Considered:** to apply hDRQN to multi-agent systems where the environment is partially observable and the task is hierarchical



(u) Multiple taxi co-operate to pick up and take off passengers



(v) Half Field Offense (A team of robots co-operates to score under the defense of another team)



(w) Multiple robots do a hierarchical tasks in a factory

Figure: Some hierarchical multi-agent domains

References |

- [1] P.-L. Bacon, J. Harb, and D. Precup, “*The option-critic architecture*,” in Proc. AAAI, 2017, pp. 1726–1734.
- [2] A. S. Vezhnevets et al. (2017). “*Feudal networks for hierarchical reinforcement learning*.” [Online]. Available: <https://arxiv.org/abs/1703.01161>
- [3] T. D. Kulkarni, K. R. Narasimhan, A. Saeedi, and J. B. Tenenbaum, “*Hierarchical deep reinforcement learning: Integrating temporal abstraction and intrinsic motivation*” in Proc. Adv. Neural Inf. Process. Syst., 2016, pp. 3675–3683
- [4] C.-C. Chiu and V.-W. Soo, “*Subgoal identifications in reinforcement learning: A survey*,” in Advances in Reinforcement Learning. Rijeka, Croatia: InTech, 2011.

- [5] M. Stolle, “*Automated discovery of options in reinforcement learning,*” Ph.D. dissertation, School Comput. Sci., McGill Univ., Montreal, QC, Canada, 2004.
- [6] M. Hausknecht and P. Stone, “*Deep recurrent Q-learning for partially observable MDPs,*” in Proc. AAAI Fall Symp. Ser., 2015.
- [7] Z. Wang, T. Schaul, M. Hessel, H. van Hasselt, M. Lanctot, and N. de Freitas. (2015). “*Dueling network architectures for deep reinforcement learning.*” [Online]. Available: <https://arxiv.org/abs/1511.06581>
- [8] H. Van Hasselt, A. Guez, and D. Silver, “*Deep reinforcement learning with double Q-learning,*” in Proc. AAAI, vol. 2, 2016, p. 5

Thank You!